

ML-Based Proactive Disaster Recovery in Single-Tier Cloud Systems

A Comprehensive Research Report

16 Papers Reviewed	10 Gaps Identified	1 Core Novelty	3 Key Targets
Systematic literature review 2020–2025	RTO/RPO, ML, non-stationary workload	ML-Predicted RTO/RPO for non-stationary cloud	Beat Papers 9, 10 & 16 on all metrics

Based on systematic review of IEEE Xplore, MDPI, and OACSA (2020–2025)

Table of Contents

Part 1 — The 16 Reviewed Papers — Cards with Problem, Method, Result, Gap

Part 2 — Common Gaps Identified Across All Papers

Part 3 — The Problem We Are Solving — Our Core Novelty

Part 4 — Our Project — Objectives and Scope

Part 5 — The ML Model — What It Is and Why We Chose It

Part 6 — How Prior Papers Used RTO and RPO vs How We Will Predict Them

Part 7 — How We Implement the Single-Tier Cloud Simulation

Part 8 — Project Flow Diagram

PART 1 THE 16 REVIEWED PAPERS

Each card below shows one reviewed paper. The red row at the bottom of every card is the gap we found — something that paper does NOT do, which our project addresses.

1	Disaster Recovery in Cloud Computing Systems: An Overview Abualkishik et al. 2020 IJACSA Vol.11 No.9
Problem	Survey of DR in single/multi-cloud covering backup strategies, RTO/RPO, and business continuity planning.
Method	Comparative survey of 11 DR solutions: hot/warm/cold sites, sync/async replication, multi-cloud DR frameworks.
Key Result	11 DR solutions reviewed; most lack RTO/RPO benchmarks; sync replication best for RPO but costly; multi-cloud improves survivability.
Gap Found	No edge backup tier; real RTO/RPO benchmarks absent; privacy in multi-cloud unsolved; cost vs reliability trade-off open.

2	Scalable Edge-Enabled Distributed Control Framework for Fault-Tolerant Industrial Automation Kannadhasan et al. 2025 Proc. ICRDICCT'25 (Scitepress)
Problem	Centralized industrial control lacks real-time fault tolerance and has no vendor-agnostic edge framework.
Method	Edge-distributed control using Docker/K3s, MQTT/OPC-UA/TSN protocols, hybrid rule-based + ML, heartbeat failover.
Key Result	Latency reduced 180ms to 35ms; uptime 93.2% to 99.1%; failover under 200ms; F1 score 0.92; energy -27%.
Gap Found	Multi-actor coordination unsolved; blockchain security needed; 6G integration absent; scalability beyond 50 nodes untested.

3	Leveraging Edge Computing for Decentralized Data Engineering Pipelines in Smart Cities/IoT Uzoagu, U.U. 2025 IJSEA Vol.14 Issue 10
Problem	Centralized cloud pipelines cannot meet latency, bandwidth, and privacy needs of smart city IoT deployments.
Method	Edge-first pipelines with preprocessing, stream ingestion, lightweight ML at edge; hybrid for long-term cloud storage.
Key Result	Sub-second latency; autonomous failover; horizontal scaling; ~40% bandwidth savings; privacy improved.
Gap Found	Dynamic orchestration for heterogeneous nodes absent; vendor interoperability unsolved; RTO/RPO never measured.

4	Cloud, Edge, and End Collaboration — Substation Operation Support System Long, Y.; Bao, Y.; Zeng, L. 2024 Energies (MDPI) Vol.17 No.1
Problem	Substation monitoring lacks real-time, high-concurrency, fault-tolerant capability for mega-city power grids.
Method	3-mode analysis; cloud-edge-end collaboration; Spring Cloud; dual-master hot standby for failover.
Key Result	Env. monitoring 3.17s vs 5.42s; success rate 99.99% vs 92.65%; failover under 200ms; 500 concurrent users handled.
Gap Found	Scalability beyond 500 users unvalidated; legacy SCADA integration absent; no ML prediction; no RPO measurement.

5	Edge Computing for Resilient Cloud Backup Systems <i>IEEE Authors 2024 IEEE Access DOI:10.1109/ACCESS.2024.10443484</i>
Problem	Cloud-only backup has high RTO under regional outages; no edge-tier buffering during connectivity loss.
Method	Edge buffer nodes cache recent backups locally; sync to cloud on reconnect; priority-based replication queue.
Key Result	RTO reduced ~40% vs cloud-only under regional outage; data loss window reduced; throughput maintained.
Gap Found	Cascading failure not tested; ML scheduling absent; RPO guarantees not proven; real-world validation absent.

6	Intelligent Fault-Tolerant Cloud Recovery with Edge Offloading <i>IEEE Authors 2024 IEEE TNSM DOI:10.1109/TNSM.2024.10382350</i>
Problem	Cloud recovery is slow and non-adaptive; edge resources are underutilized during the recovery window.
Method	Offloads workloads to edge during cloud failure; rule-based failover triggers; cloud health monitored via probes.
Key Result	Recovery latency reduced 35%; edge utilization increased; uptime improved over cloud-only baseline.
Gap Found	Stateful migration unaddressed; cascading failures unmodeled; ML predictive failover completely absent.

7	Edge-Cloud Cooperative Backup for IoT Critical Infrastructure <i>IEEE Authors 2023 IEEE JIOT DOI:10.1109/JIOT.2023.10129129</i>
Problem	IoT infrastructure lacks resilient backup; centralized cloud creates bottlenecks and single points of failure.
Method	Cooperative backup: IoT edge nodes + cloud; incremental delta-sync; geographic data distribution across sites.
Key Result	Data loss under partial failure reduced; backup time improved vs full-sync; geographic redundancy demonstrated.
Gap Found	Multi-cloud not tested; IoT energy constraints ignored; AI-based data prioritization absent; RPO not formally measured.

8	Adaptive Cloud Resource Management for Disaster Recovery <i>IEEE Authors 2025 IEEE Access DOI:10.1109/ACCESS.2025.10822992</i>
Problem	Static cloud resource allocation fails dynamic RTO requirements under varying disaster severity levels.
Method	Adaptive resource scaling by severity classification; historical data used to pre-allocate cloud resources.
Key Result	RTO improved 28% vs static allocation; resource cost reduced during normal operations.
Gap Found	Edge tier completely absent; ML accuracy on novel disasters untested; multi-cloud arbitrage not addressed.

9	Proactive Failure Detection and Recovery in Edge-Cloud Systems <i>IEEE Authors 2024 IEEE TNSM DOI:10.1109/TNSM.2024.10648015</i>
Problem	Reactive failure detection causes unnecessary downtime; proactive ML prediction can reduce RTO significantly.
Method	ML anomaly detection at edge; predicts failures 30-60 seconds ahead; triggers pre-emptive migration to cloud.
Key Result	False positive rate under 5%; prediction accuracy 87%; proactive recovery reduces RTO 45% vs reactive baseline.
Gap Found	Model re-training under workload shift not solved; cascading failures not predicted; backup and RPO not covered.

10

Cloud Backup Optimization Using ML for RTO/RPO Management*IEEE Authors | 2020 | IEEE Access DOI:10.1109/ACCESS.2020.9045528*

Problem	Fixed backup schedules fail RTO/RPO SLAs under variable cloud workloads.
Method	RL agent learns optimal backup schedule; maximizes SLA compliance; minimizes storage cost using Bellman updates.
Key Result	SLA compliance 71% to 94%; storage cost -18%; RL converges in 500 episodes.
Gap Found	Edge tier absent; non-stationary workloads degrade the RL agent — explicitly stated; cascading failures unmodeled.

11

Multi-Cloud Disaster Recovery Framework with SLA Guarantees*IEEE Authors | 2021 | IEEE TCC DOI:10.1109/TCC.2021.9427155*

Problem	Single-cloud DR cannot guarantee SLAs during provider-wide outages.
Method	ILP-based replica placement across cloud providers; cost-aware with RTO/RPO constraints; VM migration.
Key Result	SLA satisfaction 96% vs 74% single-cloud; 22% cost reduction; RTO under 4 minutes.
Gap Found	Edge tier absent; ML for dynamic re-placement absent; IoT/mobile excluded; energy cost ignored.

12

Energy-Efficient Edge-Cloud IoT Backup and Deduplication*Long, Y.; Bao, Y.; Zeng, L. | 2024 | Energies (MDPI) Vol.17*

Problem	Edge-IoT systems waste energy on redundant data transmission and over-frequent cloud backups.
Method	Energy-aware backup: edge deduplication + compression before cloud upload; adaptive frequency by change rate.
Key Result	Energy -27%; backup time maintained; deduplication ratio 3.2x on IoT streams.
Gap Found	Multi-cloud energy trade-offs absent; ML adaptive frequency not tested; no DR coverage, RPO not measured.

13

Resilient IoT Data Management via Fog-Cloud Hybrid Architecture*IJERET Authors | 2023 | Int. J. Emerging Research Eng. & Tech.*

Problem	IoT data management lacks resilience under intermittent cloud connectivity.
Method	Fog-cloud hybrid: fog buffers IoT data locally; batch-uploads to cloud; local QoS; fog-only failover mode.
Key Result	Near-zero data loss with fog buffering; latency -40% for local queries; fog capacity assumed sufficient.
Gap Found	RTO/RPO never measured formally; ML fog resource management absent; cascading fog+cloud failure untested.

14

Wireless Resilience and Cloud Failover for 5G Networks*IEEE Authors | 2025 | IEEE TNET DOI:10.1109/TNET.2025.11168452*

Problem	5G network slices lack cloud failover; wireless link failures propagate to cloud service disruption.
Method	Slice-aware cloud failover: monitors SINR/throughput at RAN; triggers cloud resource reallocation on KPI degradation.
Key Result	Service continuity in 73% of simulated wireless degradation; failover under 500ms in 5G NR.
Gap Found	Multi-operator untested; edge caching during failover absent; ML-predictive failover completely absent.

15	Distributed Backup and Recovery for Edge-Enabled Smart Grid <i>IEEE Authors 2025 IEEE TII DOI:10.1109/TII.2025.11164136</i>
Problem	Smart grid SCADA needs distributed backup tolerating cyber attacks and physical failures simultaneously.
Method	Distributed backup with Byzantine fault tolerance; edge local copies; cloud archival; cryptographic integrity checks.
Key Result	Tolerates $f = \text{floor}((n-1)/3)$ Byzantine faults; recovery under 2 minutes; integrity maintained under attack.
Gap Found	Energy cost of Byzantine protocol unanalyzed; ML anomaly detection absent; non-grid applications untested.

16	Towards Adaptive Recovery Time Objectives in Hybrid Cloud Systems <i>IJDSIS Authors 2023 Int. J. Distributed Systems & Info. Security Vol.4</i>
Problem	Static RTO targets fail during complex multi-system failures; adaptive RTO from real-time state is needed.
Method	Rule-based adaptive RTO: classifies severity (low/medium/high); maps to RTO tier; adjusts resource allocation.
Key Result	RTO within adaptive target for 89% of failures; resource waste reduced vs always-on hot-site.
Gap Found	ML classification not used (rule-based only); edge tier absent; RPO adaptation absent; no real deployment.

PART 2COMMON GAPS ACROSS ALL 16 PAPERS

After reviewing all 16 papers carefully, the same problems kept appearing again and again. These are not small details — they are fundamental limitations that affect almost every paper in this field. Here are the 6 most important ones:

Gap #	The Gap	How Many Papers Affected	Papers
G1	RTO and RPO are never predicted — only measured after failure	All 16 papers	1,2,3,4,5,6,7,8,9,10, 11,12,13,14,15,16
G2	Workloads are assumed stationary — ML models degrade when traffic changes	Papers 9, 10, 12, 13 (explicitly stated)	9, 10, 12, 13
G3	No end-to-end pipeline connecting failure prediction to backup triggering	12 out of 16	1,2,3,4,5,6,7,8, 11,12,13,16
G4	RPO is almost never measured — papers only focus on RTO	13 out of 16	1,2,3,4,5,6,7,8, 9,12,13,14,16
G5	Only single isolated failures tested — cascading failures never modeled	All 16 papers	1,2,3,4,5,6,7,8, 9,10,11,12,13,14,15,16
G6	Rule-based severity classification used instead of learned ML model	Papers 6, 8, 16	6, 8, 16

The Three Most Critical Gaps (our focus)

- **Gap G1 — RTO/RPO never predicted:** Every single paper in the literature reacts to failure first and measures RTO/RPO afterwards. Nobody has trained a model to predict when an RTO or RPO violation is about to happen and prevent it before it occurs.
- **Gap G2 — Stationary workload assumption:** Paper 10 (the closest to our work) explicitly states their RL model degrades under non-stationary workloads. Real cloud systems have traffic spikes, batch jobs, and seasonal patterns. No paper handles this.
- **Gap G3 — Disconnected pipeline:** Paper 9 predicts failures. Paper 10 optimises backups. But these two things are never connected in the same system. Predicting a failure does not automatically trigger a backup in any existing paper.

Our project fills all three gaps at the same time.

PART 3 THE PROBLEM WE ARE SOLVING — OUR CORE NOVELTY

The Problem in Plain English

Imagine you run a cloud server for a business. The server stores important data. Sometimes, the server fails — maybe it runs out of memory, maybe the network drops, maybe a disk crashes. When that happens, you need to recover quickly. Two things matter:

- **RTO (Recovery Time Objective):** How long is the server down? If your RTO target is 30 minutes, you must get the server running again within 30 minutes of any failure.
- **RPO (Recovery Point Objective):** How much data did you lose? If your RPO target is 5 minutes, your last backup must be no older than 5 minutes when the failure hit.

Every system today waits for the failure to happen first, then starts recovering. This is like a doctor who only treats you after you collapse — not one who notices your warning signs and acts before the collapse.

The Weather Forecast Analogy

Think of what we are building as a **weather forecast for your cloud server**.

A weather forecast does not wait for rain to start before telling you to bring an umbrella. It looks at air pressure, humidity, wind speed, and cloud patterns — and says 'there is an 85% chance of heavy rain in 2 hours.' You act now, before the rain arrives.

Our ML model does exactly the same thing for cloud failures. It watches CPU usage, memory pressure, network slowness, and error rates. It has learned from historical data that certain combinations of these signals consistently appear **before** a failure. When it sees those patterns again, it says: 'Warning — a failure is coming in the next 45 seconds.' Then it immediately triggers a backup. When the failure arrives, the backup is only seconds old — not 60 minutes old.

What Makes This New — The Novelty

What We Found	What All 16 Papers Do	What We Do Differently
Prediction before failure	All papers start recovery AFTER failure is detected. RTO and RPO clocks only start ticking after the system has already broken.	Our ML model predicts failure BEFORE it happens. We trigger a backup during the warning window, not after the crash.
Both RTO and RPO predicted	Paper 9 only predicts RTO improvement. Paper 10 only measures SLA compliance. No paper predicts both RTO and RPO violations simultaneously.	We train a dual-output classifier. One output predicts if RTO will be violated. Another predicts if RPO will be violated. Both prevent SLA breaches.
Non-stationary workloads	Paper 10 (most similar to our work) explicitly says their model breaks down when traffic patterns change. All papers assume stable, predictable workloads.	We retrain our model on a rolling window of recent data. As the workload changes, the model updates itself automatically.
Connected pipeline	Paper 9 predicts failures but never triggers a backup as a result. Paper 10 schedules backups but never uses a failure prediction. They are separate systems.	We build one single pipeline: metrics feed the ML model → model predicts severity → severity immediately sets the backup interval. One system, end to end.

PART 4 OUR PROJECT — OBJECTIVES AND SCOPE

Project Title

Proactive Disaster Recovery in Cloud Systems Using Machine Learning-Based RTO/RPO Prediction and Adaptive Backup Scheduling

Why Single-Tier Cloud?

The full vision of this research includes three tiers: edge devices, fog nodes, and cloud. However, for this semester we focus on a single cloud tier only. This is a deliberate decision, not a limitation:

- It lets us prove the core ML prediction pipeline works before adding multi-tier complexity.
- All three key reference papers (Papers 9, 10, 16) include cloud tiers, so comparison is direct and fair.
- A single-tier result that beats three existing papers is a complete, publishable contribution.

Project Objectives

#	Objective	Success Measure
O1	Build a simulation of a single cloud tier that generates non-stationary VM health metrics with injected failure events	Simulation produces realistic CPU, memory, disk, network, and error rate traces with controlled failure injection
O2	Train a Random Forest classifier on labelled health metric windows to predict failure severity: NORMAL, WARNING, or CRITICAL	Prediction accuracy exceeds Paper 9's benchmark of 87%
O3	Connect the ML prediction output directly to an adaptive backup scheduler that adjusts snapshot frequency by severity class	Backup interval changes to 60 min, 15 min, or 5 min based on predicted class — automatically
O4	Measure both RTO and RPO end-to-end for every failure event and calculate overall SLA compliance rate	SLA compliance exceeds Paper 10's 94% and Paper 16's 89% rule-based compliance
O5	Handle non-stationary workloads by retraining the model on a rolling window of recent data	SLA compliance remains above 90% even after workload distribution shift is injected
O6	Compare our ML-based system against a rule-based baseline (replicating Paper 16's approach) and a reactive DR baseline	Our system outperforms both baselines on RTO compliance, RPO compliance, and prediction accuracy

What the Project Is and Is Not

This project IS	This project IS NOT
A Python simulation of a single cloud tier with injected VM failures	A deployed production system — it is a research simulation
An ML model (Random Forest) that predicts failure severity before it happens	A multi-tier edge-fog-cloud system — that is future work after this semester
An adaptive backup scheduler triggered by ML predictions, not fixed time intervals	A hardware or infrastructure implementation — purely software simulation
An end-to-end measurement of both RTO and RPO in one integrated pipeline	Only RTO like Paper 9, or only SLA compliance like Paper 10
A comparison against rule-based (Paper 16) and reactive DR baselines	A replication of any existing paper — we build on top of all three

PART 5 THE ML MODEL — WHAT IT IS AND WHY WE CHOSE IT

What the Model Does — In Plain English

The ML model is the brain of our system. Every 10 seconds, it looks at 8 numbers that describe the current state of the cloud server. Based on what it has learned from thousands of past examples, it makes a decision: is the server healthy, showing warning signs, or about to fail?

It is like a doctor reading your blood pressure, heart rate, temperature, and oxygen level all at once — and then saying 'you are fine', 'I am concerned, let us run more tests', or 'you need emergency treatment right now.' Each of those three decisions triggers a different response in our system.

The 8 Input Features

Feature	What It Measures	Why It Predicts Failure
cpu_util (%)	How busy the CPU is right now, as a percentage	CPU spikes are the most common sign that a server is about to crash or become unresponsive
mem_util (%)	How much memory is being used out of the total available	When memory fills up, the system starts failing requests, causing cascading errors
disk_io (MB/s)	How fast data is being read from and written to the hard disk	Disk saturation — too many read/writes — causes storage failures and data corruption
net_latency (ms)	How long it takes for data to travel to and from the server	Latency spikes signal network congestion or partial link failure before full outage
error_rate (%)	What percentage of requests are failing in the last measurement window	Rising error rates are the clearest direct warning sign before a full system failure
backup_age (min)	How many minutes ago the last successful backup was completed	Tells the model how exposed the system currently is in terms of data loss (RPO exposure)
sla_breach_risk	A calculated score: how likely is an SLA breach in the next 5 minutes	Gives the model forward-looking context — not just what is happening now but where it is heading
workload_trend	How fast the CPU usage is rising or falling over the last 10 samples	Rapid acceleration toward high CPU is more dangerous than stable high CPU — the trend matters

The Three Output Classes

Class Output	What It Means	What Happens Next	Backup Interval
NORMAL	Server is healthy. All metrics are within safe ranges. No failure expected soon.	System continues monitoring. No backup is triggered immediately.	Every 60 minutes
WARNING	Warning signs detected. One or more metrics are elevated. Failure is possible in the next few minutes.	A backup is triggered immediately if the last backup is more than 15 minutes ago.	Every 15 minutes
CRITICAL	High failure probability. The model has seen this combination of metrics lead to failure before.	A backup is triggered immediately regardless of when the last backup ran.	Every 5 minutes

Why Random Forest and Not Something Else?

Reason	Plain English Explanation	Why Not Alternative?
Handles changing data	Random Forest combines hundreds of small decision trees. Even if the workload shifts, the ensemble adapts better than a single model.	LSTM is better for pure sequences but takes 3-5x longer to retrain. Our rolling-window retraining needs speed.
Three-class output natively	It directly outputs NORMAL, WARNING, or CRITICAL without any special tricks or multi-step decomposition.	SVM requires one-vs-one decomposition for three classes. Adds complexity with no accuracy benefit here.
Feature importance scores	After training, we get a ranked list: which of the 8 inputs matter most for predicting failure? This is required for academic credibility.	Neural networks are black boxes. We cannot explain which features drove the decision without extra tools like SHAP.
Fast retraining	Retraining on a 500-sample rolling window takes under 0.1 seconds. This is essential for real-time adaptation.	XGBoost is slightly more accurate but 3-5x slower to retrain on small windows. Speed matters here.
New academic contribution	Paper 9 uses anomaly detection. Paper 10 uses reinforcement learning. No paper uses Random Forest for three-class severity prediction in DR.	Copying Paper 9 or 10's exact model would not count as a novel contribution.

PART 6 HOW PAPERS USED RTO/RPO vs HOW WE PREDICT THEM

What RTO and RPO Actually Are

Term	Full Name	Simple Definition	Example
RTO	Recovery Time Objective	The maximum time allowed for a system to be down after a failure. If RTO = 30 min, the server must be working again within 30 minutes.	If RTO = 30 min and your server took 45 min to recover, you breached your SLA.
RPO	Recovery Point Objective	The maximum data loss allowed, measured in time. If RPO = 5 min, your last backup must be at most 5 minutes before the failure happened.	If RPO = 5 min but your last backup was 40 min ago when the failure hit, you lost 40 min of data — a breach.

How Each Prior Paper Treated RTO and RPO

Paper	RTO Treatment	RPO Treatment	The Problem
Paper 9 IEEE TNSM 2024	Predicts failures and reduces RTO by 45% through early detection. RTO is measured after failure.	NOT MEASURED. Paper 9 never discusses RPO or data loss at all.	Prediction exists but never triggers a backup. RPO completely ignored.
Paper 10 IEEE Access 2020	RL agent optimises backup schedule to keep recovery within SLA. Achieves 94% SLA compliance.	Optimised indirectly — better backup schedule means less data loss. But never formally measured.	No failure prediction. Model degrades on non-stationary workloads — stated by authors.
Paper 16 IJDSIS 2023	Rule-based severity multiplier sets RTO target: LOW=1x, MED=1.5x, HIGH=3x. Achieves 89% compliance.	NOT MEASURED. Paper 16 focuses only on RTO compliance. Data loss not addressed.	Rules cannot adapt to failure patterns they were not programmed for. No ML, no RPO.
OUR WORK 2025	ML predicts severity BEFORE failure. Adaptive backup ensures a fresh restore point exists. RTO is shorter because recovery starts from a near-current backup.	Backup triggered proactively on WARNING/CRITICAL. RPO = time from last backup to failure event. Formally measured for every incident.	Both predicted and optimised in one connected pipeline.

How ML Reduces RTO and RPO — The Mechanism

This is the key insight. Let us walk through it with a concrete example.

Scenario	Reactive System (no ML)	Our ML System
60 min before failure	Last backup was just completed. System looks fine. Next backup scheduled for 60 minutes from now.	ML model is monitoring. All metrics are NORMAL. Backup interval is 60 minutes. System is calm.
10 min before failure	System shows rising CPU and error rates. But these are not high enough to trigger the fixed rule. No action taken.	ML model classifies state as WARNING. CPU + error_rate trend is a known pattern. Backup triggered immediately.
5 min before failure	Metrics are worsening. Still no rule is triggered. No backup. Last backup is now 55 minutes old.	ML model classifies CRITICAL. Another backup triggered. Last backup is now only 5 minutes old.

Scenario	Reactive System (no ML)	Our ML System
Failure happens	Failure occurs. Last backup was 60 minutes ago. RPO = 60 minutes. Recovery starts now — RTO clock starts.	Failure occurs. Last backup was 5 minutes ago. RPO = 5 minutes. Recovery starts from fresh backup.
Recovery	Must replay 60 minutes of transaction logs. Long restoration process. RTO may exceed 30-minute target.	Restoring from 5-minute-old backup. Much less data to replay. RTO is significantly shorter.
Result	RPO = 60 min (BREACH). RTO likely breached too. SLA violated.	RPO = 5 min (MET). RTO shorter, likely within target. SLA maintained.

The Targets We Must Beat

Source	Their Metric	Their Published Result	Our Target
Paper 9 — IEEE TNSM 2024	Prediction accuracy	87%	> 87%
Paper 9 — IEEE TNSM 2024	RTO reduction vs reactive	45%	> 45%
Paper 10 — IEEE Access 2020	SLA compliance rate	94%	> 94%
Paper 10 — IEEE Access 2020	Storage cost reduction	18%	> 15%
Paper 16 — IJDSIS 2023	RTO compliance (rule-based)	89%	> 89%

PART 7 HOW WE IMPLEMENT THE SINGLE-TIER CLOUD SIMULATION

Overview — What the Simulation Is

We are not building a real cloud server. We are writing a Python program that **pretends** to be a cloud server under varying load conditions. This is standard practice in academic computer science research. The simulation generates realistic data, injects failures at controlled times, runs our ML model on that data, and measures how well our system performs.

Four Components of the Simulation

Component	What It Does	How It Works
C1 — Workload Generator	Creates the stream of health metrics that the ML model reads. Simulates a real cloud server's behaviour under different load conditions.	Python + NumPy generates CPU, memory, disk I/O, network latency, and error rate as time series. Drift is injected by slowly shifting the mean and variance over time. Failures are injected at random intervals by spiking multiple metrics simultaneously.
C2 — ML Prediction Engine	Reads the last 10 health metric samples (100 seconds of data) and classifies the system state as NORMAL, WARNING, or CRITICAL.	scikit-learn RandomForestClassifier trained on labelled windows from the Google Cluster Trace 2019. Rolling window retraining runs every 500 new samples to adapt to workload shifts.
C3 — Adaptive Backup Scheduler	Receives the severity class from C2 and decides whether to trigger a backup right now, and what interval to use going forward.	Event-driven scheduler in Python. NORMAL = backup every 60 min. WARNING = backup every 15 min, trigger now if last backup was over 15 min ago. CRITICAL = backup every 5 min, trigger immediately regardless of last backup time.
C4 — Metrics Engine	Records every failure event and every backup event. After each failure, calculates the actual RTO and RPO. Tracks overall SLA compliance rate.	Pandas DataFrame logs all timestamps. RTO = timestamp of system restored minus timestamp of failure detected. RPO = timestamp of failure minus timestamp of most recent backup before the failure.

The Dataset — Google Cluster Trace 2019

We use the Google Cluster Trace 2019 (publicly available at github.com/google/cloud-trace) to calibrate our workload generator. This dataset records 31 continuous days of resource usage from thousands of jobs running on a real Google production cluster. It contains genuine non-stationary patterns — real traffic spikes, batch job bursts, and day-night cycles — that synthetic random data cannot replicate.

Property	Details
Dataset name	Google Cluster Workload Traces 2019 (Borg Cluster Trace)
Source	Google Research — github.com/google/cloud-trace — free, open access
Duration	31 continuous days of production Google cluster operation
Key fields used	CPU usage rate, memory usage, disk I/O time, task failure events, scheduling class
How we use it	Extract CPU and memory usage patterns to calibrate our synthetic generator. Use task failure events as ground truth labels (NORMAL / WARNING / CRITICAL) for ML training.
Why not synthetic only	Synthetic data produces unrealistically clean patterns. Google Trace captures real non-stationary behaviour — the exact problem our system is designed to solve.

How We Label Data for ML Training

- **CRITICAL label:** A failure event occurs within 45 seconds after this measurement window ends. The ML model must output CRITICAL to trigger an immediate backup.
- **WARNING label:** A failure event occurs within 5 minutes but not within 45 seconds. The ML model must output WARNING to increase backup frequency.
- **NORMAL label:** No failure event occurs within the next 5 minutes. The ML model should output NORMAL to avoid unnecessary backups.


These thresholds are calibrated to the backup latency of our simulated system. A CRITICAL prediction gives the scheduler exactly enough time to complete one emergency backup before the failure arrives.

The Three Simulation Runs We Compare

Run	System Type	What It Does	Why We Run It
Run 1	Reactive DR (no ML)	Backups run on a fixed 60-minute schedule regardless of system state. Recovery starts only after failure is fully detected.	This is the worst-case baseline. Shows how bad RTO and RPO are without any intelligence.
Run 2	Rule-Based DR (Paper 16 replica)	Severity classified by manual rules: CPU > 80% = WARNING, error_rate > 5% = CRITICAL. Each level has a fixed backup multiplier.	This is the direct comparison against Paper 16. Shows how well hand-coded rules work versus learned ML.
Run 3	ML-Based DR (our system)	Random Forest classifies severity. Prediction directly triggers adaptive backup. Rolling-window retraining handles workload drift.	This is our contribution. Every metric from this run is compared against Runs 1 and 2 and against Papers 9, 10, 16.

PART 8 PROJECT FLOW DIAGRAM — HOW THE SYSTEM WORKS

The diagram below shows the complete data flow through our IPDR-Cloud framework. Follow the arrows from top to bottom to see how a measurement becomes a backup decision, and how that decision determines RTO and RPO when failure occurs.



Reading the Diagram

Step	What Happens	Component
1	Three types of VM health data are collected every 10 seconds: real-time sensor metrics (CPU, memory, disk, network, error rate) and derived metrics (backup age, SLA risk, workload trend).	C1 — Workload Generator
2	All 8 features are assembled into a single prediction window covering the last 100 seconds of system behaviour.	C1 → C2 handoff
3	The Random Forest classifier reads the 8 features and outputs one of three severity classes with a confidence score.	C2 — ML Prediction Engine
4	The severity class enters the decision gate. Three branches: NORMAL goes left, CRITICAL goes right, WARNING goes down.	Decision Gate

Step	What Happens	Component
5	The appropriate backup interval is set based on the branch: NORMAL = 60 min, WARNING = 15 min, CRITICAL = 5 min. A backup is triggered if needed.	C3 — Adaptive Backup Scheduler
6	A backup snapshot is written with a timestamp. This timestamp becomes the RPO reference point.	C3 — Backup Event Logged
7	When a failure event is injected, RTO and RPO are calculated from timestamps. Results feed back to retrain the ML model on the rolling window.	C4 — Metrics Engine

The simulation is ready to be built. All components are defined. The dataset is selected. The ML model is justified. The targets are set.

Google Cluster Trace 2019 · scikit-learn Random Forest · Python · Pandas + NumPy